

Simple, Fast, and Practical Non-Blocking and Blocking Concurrent Queue Algorithms

Pseudocode from [article](#) of the above name in *PODC96* (with two typos corrected), by [Maged M. Michael](#) and [Michael L. Scott](#). Corrected version also appeared in *JPDC*, 1998.

The [non-blocking concurrent queue algorithm](#) performs well on dedicated as well as multiprogrammed multiprocessors with and without contention. The algorithm requires a universal atomic primitive, *CAS* or *LL/SC*. It depends for memory management on a type-preserving allocator that never reuses a queue node as a different type of object, and never returns memory to the operating system. If this is unacceptable in a given context, the code can be modified to incorporate [hazard pointers](#), [epoch-based reclamation](#), or [interval-based reclamation](#).

The [two-lock concurrent queue algorithm](#) performs well on dedicated multiprocessors under high contention. Useful for multiprocessors without a universal atomic primitive.

Non-Blocking Concurrent Queue Algorithm

```
structure pointer_t {ptr: pointer to node_t, count: unsigned integer}
structure node_t {value: data type, next: pointer_t}
structure queue_t {Head: pointer_t, Tail: pointer_t}

initialize(Q: pointer to queue_t)
    node = new_node()           // Allocate a free node
    node->next.ptr = NULL        // Make it the only node in the linked list
    Q->Head.ptr = Q->Tail.ptr = node // Both Head and Tail point to it

enqueue(Q: pointer to queue_t, value: data type)
    E1: node = new_node()       // Allocate a new node from the free list
    E2: node->value = value      // Copy enqueued value into node
    E3: node->next.ptr = NULL    // Set next pointer of node to NULL
    E4: loop                    // Keep trying until Enqueue is done
    E5:     tail = Q->Tail       // Read Tail.ptr and Tail.count together
    E6:     next = tail.ptr->next // Read next ptr and count fields together
    E7:     if tail == Q->Tail // Are tail and next consistent?
        // Was Tail pointing to the last node?
    E8:     if next.ptr == NULL
        // Try to link node at the end of the linked list
    E9:     if CAS(&tail.ptr->next, next, <node, next.count+1>)
    E10:         break // Enqueue is done. Exit loop
    E11:     endif
    E12:     else // Tail was not pointing to the last node
        // Try to swing Tail to the next node
    E13:         CAS(&Q->Tail, tail, <next.ptr, tail.count+1>)
    E14:     endif
    E15:     endif
    E16: endloop
        // Enqueue is done. Try to swing Tail to the inserted node
    E17: CAS(&Q->Tail, tail, <node, tail.count+1>)

dequeue(Q: pointer to queue_t, pvalue: pointer to data type): boolean
```

```

D1:  loop                                // Keep trying until Dequeue is done
D2:    head = Q->Head                    // Read Head
D3:    tail = Q->Tail                     // Read Tail
D4:    next = head.ptr->next              // Read Head.ptr->next
D5:    if head == Q->Head                  // Are head, tail, and next consistent?
D6:        if head.ptr == tail.ptr        // Is queue empty or Tail falling behind?
D7:            if next.ptr == NULL        // Is queue empty?
D8:                return FALSE           // Queue is empty, couldn't dequeue
D9:        endif
D10:       // Tail is falling behind. Try to advance it
D11:       CAS(&Q->Tail, tail, <next.ptr, tail.count+1>)
D12:     else                             // No need to deal with Tail
D13:       // Read value before CAS
D14:       // Otherwise, another dequeue might free the next node
D15:       *pvalue = next.ptr->value
D16:       // Try to swing Head to the next node
D17:       if CAS(&Q->Head, head, <next.ptr, head.count+1>)
D18:           break                       // Dequeue is done. Exit loop
D19:       endif
D20:     endif
D21:   endloop
D22:   free(head.ptr)                     // It is safe now to free the old node
D23:   return TRUE                         // Queue was not empty, dequeue succeeded

```

Two-Lock Concurrent Queue Algorithm

```

structure node_t {value: data type, next: pointer to node_t}
structure queue_t {Head: pointer to node_t, Tail: pointer to node_t,
                  H_lock: lock type, T_lock: lock type}

initialize(Q: pointer to queue_t)
    node = new_node()                    // Allocate a free node
    node->next = NULL                     // Make it the only node in the linked list
    Q->Head = Q->Tail = node              // Both Head and Tail point to it
    Q->H_lock = Q->T_lock = FREE          // Locks are initially free

enqueue(Q: pointer to queue_t, value: data type)
    node = new_node()                    // Allocate a new node from the free list
    node->value = value                    // Copy enqueued value into node
    node->next = NULL                     // Set next pointer of node to NULL
    lock(&Q->T_lock)                       // Acquire T_lock in order to access Tail
    Q->Tail->next = node                  // Link node at the end of the linked list
    Q->Tail = node                        // Swing Tail to node
    unlock(&Q->T_lock)                     // Release T_lock

dequeue(Q: pointer to queue_t, pvalue: pointer to data type): boolean
    lock(&Q->H_lock)                       // Acquire H_lock in order to access Head
    node = Q->Head                         // Read Head
    new_head = node->next                  // Read next pointer
    if new_head == NULL                    // Is queue empty?
        unlock(&Q->H_lock)                 // Release H_lock before return
        return FALSE                       // Queue was empty
    endif
    *pvalue = new_head->value              // Queue not empty. Read value before release
    Q->Head = new_head                     // Swing Head to next node
    unlock(&Q->H_lock)                     // Release H_lock
    free(node)                             // Free node
    return TRUE                           // Queue was not empty, dequeue succeeded

```
